

SPI, DMA, & Real-time Operating Systems

February 18, 2013

Table of contents

SPI

DMA

Threads

Resource Management

Thread Synchronization


Hardware Abstraction

Serial Peripheral Interface (SPI)

- ▶ Common, 4-pin, multi-slave, full-duplex bus
- ▶ Capable of high speeds and uses a very simple hardware interface
- ▶ Master generates clock
- ▶ Both master and slave have a data output line to provide data on the clock
- ▶ Slaves selected by active-low select line (nCS, nSS) – one per slave
- ▶ Data always goes in both directions – send *dummy bytes* if nothing to send

SPI Parameters

- ▶ Clock speed – usually 1-30MHz
- ▶ Clock polarity (CPOL) – value of clock line when idle
- ▶ Clock phase (CPHA) – which edge (first or second) latches new data
 - ▶ If $CPHA = 1$, first edge reads data, second edge shifts new data to transmit



SPI Virtual-Register Protocol

- ▶ Most SPI slaves implement a virtual-register based protocol
- ▶ First byte contains address, a read/write flag, and a burst flag
- ▶ Subsequent bytes contain data to be transferred
- ▶ Burst indicates that multiple bytes will be accessed sequentially and the address should be automatically incremented

SPI Virtual-Register Protocol

Read Example

Master	Addr & R/nW & Burst			
Slave		Data1	Data2	Data3

Write Example

Master	Addr & Burst	Data1	Data2	Data3
Slave				

Direct Memory Access (DMA)

- ▶ AHB matrix in STM32F4 has both master and slave devices
- ▶ Slave – Cannot initiate transactions
 - ▶ RAM
 - ▶ Flash
 - ▶ Peripherals
- ▶ Master
 - ▶ CPU (multiple masters)
 - ▶ USB
 - ▶ Ethernet
 - ▶ *DMA1 & DMA2* (2 masters each)

DMA Motivations

- ▶ For use of serial peripherals ($<1\text{MHz}$ word rate), it is a waste of CPU time to move data
- ▶ Wait thousands of cycles for byte to be complete \rightarrow copy byte \rightarrow repeat
- ▶ \uparrow dumb
- ▶ DMA is a dumb device dedicated to this task
- ▶ It knows when and where to move data and how to raise an interrupt for the CPU when a specified amount of data has been moved
- ▶ As a bus master, it has the power to perform all the same transactions that the CPU can (including memory-to-memory transfers)

Alternate DMA Motivations

- ▶ Some devices such as USB, Ethernet, and audio may require data to be moved as fast as the CPU possibly could
- ▶ This would leave the CPU saturated and with no time to run your application
- ▶ Memory-to-memory copies can also be done efficiently with DMA if configuration time is tolerable

DMA Structure

Streams

- ▶ A stream contains all of the resources to move data from one place to another
- ▶ Up to 8 may be active at once per DMA device – Use priority-based arbitration
- ▶ Each stream gets its own interrupt vector

Channels

- ▶ Essentially a trigger source, associating a stream with a peripheral
- ▶ Allow the peripheral to indicate when it is ready for another transfer

DMA Structure

FIFOs

- ▶ 4-word buffer for packing/unpacking to optimize memory bandwidth
 - ▶ If AHB nearly saturated, preferable to move one 32-bit value rather than 4 8-bit values
- ▶ Bus transfers initiated at 1/4, 1/2, 3/4 or full
 - ▶ Burst transfers can be faster as arbitration may be held across several bytes

DMA Configuration

- ▶ Configuration is just like any other peripheral in the STM32 ecosystem
- ▶ Parameters:
 - ▶ Peripheral address – The address of the memory-mapped peripheral register to read or write
 - ▶ Memory address – The base address for the memory transfer
 - ▶ Direction
 - ▶ Request channel (Table 20/21 in FRM)
 - ▶ Number of words
 - ▶ Word size
 - ▶ Circular/double buffering
 - ▶ Priority
 - ▶ FIFO configuration (or direct mode)
 - ▶ Interrupt configuration

DMA Buffering Options

Circular Buffering

- ▶ DMA will continue when it finishes, resetting the memory address and count each time
- ▶ Use when maximum throughput is required (no overhead restarting transfer)
- ▶ Interrupts can be generated at different points in the transfer to allow time to reload/read buffer

Double Buffering (Implies Circular)

- ▶ Two base addresses are provided and DMA switches between the two
- ▶ If the two buffers are adjacent, this is the same as circular buffering

Memory-to-peripheral Transfers

- ▶ Peripheral must indicate to DMA controller (request) when new data required
- ▶ `SPI_DMACmd(SPI1, SPI_DMAReq_Tx, ENABLE);`
- ▶ Generally want memory pointer to increment, peripheral pointer constant
- ▶ Unless the target peripheral has a FIFO built in, must transfer one word at a time
- ▶ Set word size to be data size for peripheral

Peripheral-to-memory Transfers

- ▶ Peripheral must indicate to DMA controller (request) when new data available
- ▶ `SPI_DMAMCmd(SPI1, SPI_DMAREq_Rx, ENABLE);`
- ▶ Generally want memory pointer to increment, peripheral pointer constant
- ▶ May buffer multiple words in DMA FIFO
- ▶ Set word size to be data size for peripheral

Full-Duplex with DMA

- ▶ Requires two streams
- ▶ For SPI master, read-through-write creates dependency of RX on TX
 - ▶ Have to start RX DMA before TX

Real-time Operating Systems (RTOS)

An RTOS....

- ▶ ... provides basic multi-tasking support (threads)
- ▶ ... provides resource management
- ▶ ... provides thread synchronization utilities
- ▶ ... provides varying degrees of hardware abstraction

You will be using Keil's RTX with a CMSIS-RTOS API wrapper

REFERENCE MATERIAL

Anything I don't cover about this:

<http://bennahill.com/docs/cmsis/CMSIS/Documentation/RTOS/html/>

Threads

- ▶ Thread support allows multiple tasks to run concurrently
- ▶ A *scheduler* determines which should be running and for how long
 - ▶ For this course, scheduler is *round-robin*
- ▶ Threads have their own stack which allows them persistent data storage

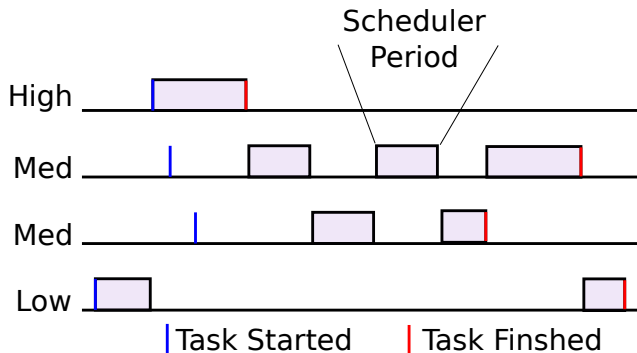
Round-Robin (RR) Scheduler

- ▶ Simple *priority-based, time-sharing, pre-emptive* scheduling scheme
- ▶ Very common among low-power embedded systems
- ▶ *Period* or *tick rate* refers to the frequency of context switches
 - ▶ Faster means more even distribution among tasks but more overhead in context switches

Schedule

- ▶ Tasks ready to execute in *ready pool*
- ▶ Highest-priority threads in ready pool are cyclically given a fixed time to execute
- ▶ Threads sleep while delaying or waiting on a resource or synchronization

RR Scheduling Example



Priority Selection

- ▶ Objective is to have *all* tasks *always* finish when they need to
- ▶ For this class, use a fixed-priority scheme
- ▶ Interrupts are at a higher priority than any thread
 - ▶ Limit time here to minimize delays in threads
- ▶ All non-real-time tasks (no deadline) minimum priority
- ▶ For others, prioritize based on how quickly it must finish

Soft Timers

Differ from “hard” timers in that they don’t use a hardware timer directly and are instead managed by scheduler ticks

- ▶ Don’t necessarily have guarantees regarding the actual time of a delay
- ▶ Generally limited to the granularity of the scheduler tick
- ▶ Upon expiration, a callback function is executed, usually at a high priority (like an interrupt)

One-time A single timer that expires only once and stops

Periodic Reloads value upon expiration and maintains periodic behavior

CMSIS-RTOS Soft-Timer API

```
typedef struct {  
    os_ptimer ptimer; // Callback function address  
                        // (takes void * parameter)  
} osTimerDef_t;  
  
typedef enum {  
    osTimerOnce = 0,  
    osTimerPeriodic = 1  
} os_timer_type;  
  
osTimerId osTimerCreate(osTimerDef_t *timer_def,  
                        os_timer_type type, void *argument);  
  
osStatus osTimerStart (osTimerId timer_id, uint32_t millisec);  
  
osStatus osTimerStop (osTimerId timer_id);
```

Thread Management

<http://bennahill.com/docs/cmsis/CMSIS/Documentation/RTOS/html/>

Creation

- ▶ Macro shortcut: `osThreadDef (name,priority,instances,stacksz);`
 - ▶ Make a thread with specified priority and create a few instances with statically allocated stacks



`osThreadId osThreadCreate(osThreadDef_t *thread_def, void *argument)`

Basic Operations

Yield Give up the allotted scheduler slot for now – will get it back

Terminate(id) Kill a thread

SetPriority(p) Assign a new priority

Resource Management

“Resources”

Generally, items that you need to control access to because multiple threads might want to use simultaneously

Could be:


- ▶ Variables or memory regions (i.e. sensitive data structures like linked lists)
- ▶ Peripherals (don't want to interrupt transmissions)
- ▶ Flow control and synchronization resources

Mutual Exclusion (Mutex)

- ▶ At most one thread may hold a mutual exclusion “lock”
- ▶ Other threads that wait will sleep until lock is released
- ▶ When a lock is released, it is given to the highest priority waiting thread
- ▶ `osMutexId osMutexCreate (osMutexDef_t *mutex_def);`
- ▶ `osStatus osMutexWait (osMutexId mutex_id, uint32_t millisec);`
- ▶ `osStatus osMutexRelease (osMutexId mutex_id);`

Priority Inversion

A basic fixed-priority scheduler has potential problems with priority inversion, where a pair of tasks run with inverted relative priorities



Note that the medium priority tasks, even without holding any shared resources, prevent the high-priority task from executing

Priority Inheritance Protocol

To fix this, most RTOS implementations use Priority Inheritance Protocol (PIP)

- ▶ When a task holds a mutex, it runs with the priority of the highest-priority waiting thread
- ▶ This effectively prioritizes the release of the lock according to those that need it
- ▶ This isn't perfect since problems can come up in nested cases, but it's easy and requires minimal prior knowledge

Semaphores

Similar to a mutex, but does not associate ownership by any thread

```
osSemaphoreId osSemaphoreCreate (osSemaphoreDef_t *semaphore_def,  
int32_t count);
```

- ▶ Associates well with countable resources
- ▶ Two basic operations:
 - ▶ Wait – Blocks as long as the value of the semaphore is 0, then decrement

```
int32_t osSemaphoreWait(osSemaphoreId semaphore_id,  
uint32_t millisec);
```
 - ▶ Release – Increment the value of the semaphore

```
osStatus osSemaphoreRelease(osSemaphoreId semaphore_id);
```

Semaphores with a maximum count of 1 are called *binary semaphores*

Mutex vs. Semaphore

- ▶ Semaphores are faster in every way – very quick locks (protecting an operation that may only take a few cycles) can take advantage of this
- ▶ Semaphores have no knowledge of thread ownership and thus can't use PIP
- ▶ Semaphores can easily associate with multiple identical functional units

Thread Synchronization

Most RTOSs provide a number of ways to safely and efficiently move data of various sizes between threads

- ▶ The names and conventions of each of these devices varies greatly across different OS implementations
- ▶ Generally: Queues (arbitrarily-sized data), events (minimal data), and something in between (usually integers/pointers)

Events (CMSIS-RTOS “Signals”)

These are minimally sized data intended purely to indicate that something has happened

- ▶ Used for timing and very simple I/O
- ▶ Implemented as a 32-bit integer where each bit represents a different application-designated signal

Basic Operations

`set(thread, flag mask)` Set a bit mask on a thread – OR its current signal flags with your provided mask

`wait(flag mask)` Wait for the specified signal flags to be set – Clear selected flags once all have been set.

Queues (CMSIS-RTOS “Mail Queue”)

Send copies of arbitrarily sized data to a mailbox

- ▶ Includes efficient (fixed block size) memory allocation pool
- ▶ Often used for high-level gatekeeping tasks (i.e. dispatching or receiving large packets of data from a peripheral, perhaps accelerometer data...)
- ▶ Queue length limited by number of buffers

Basic Operations

allocate Allocate a block from the pool and return a pointer to it so it can be filled with data – this must later be freed

put(data) Add a block to the queue

get Blocks until data available – Pointer to data returned (you should free it now)

free(data) Free a block now that it has been received

Message Queue

A simpler queue which just sends integers

- ▶ Flexible for you to design your own more efficient protocol
- ▶ No explicit allocation required but messages can be pointers
- ▶ Queue length limited by a buffer size

Basic Operations

`put (uint)` Add data to the queue

`get` Wait for data in queue and return it

Hardware Abstraction

CMSIS-RTOS provides no real hardware abstraction